

# SSN College of Engineering, Chennai

## Department of Computer Science and Engineering

**Date:** July 30, 2025

**Academic Year:** 2025–2026

**Degree:** M.Tech (5-Year Integrated) CSE

**Course Code:** ICS1512 Machine Learning Algorithms Laboratory

### Experiment 1: Working with Python packages-Numpy, Scipy, Scikit-Learn, Matplotlib

#### AIM

To explore five machine learning datasets, select and justify the most suitable algorithms for each task, implement a workflow in Python, including data pre-processing, model training, and performance evaluation, and interpret the results using appropriate metrics.

#### Libraries Used and Key Concepts

This section highlights commonly used Python libraries in machine learning and explores key functions and concepts from each.

##### 1. NumPy

**Purpose:** Numerical computing and array manipulation.

**Key Functions:**

- `np.array()`: Create arrays
- `np.mean()`, `np.std()`, `np.sum()`: Statistical computations
- `np.dot()`, `np.linalg.inv()`: Linear algebra

## 2. Pandas

**Purpose:** Data manipulation and analysis.

**Key Functions:**

- `pd.read_csv()`: Load data from CSV
- `df.head()`, `df.describe()`, `df.info()`: Data inspection
- `df.fillna()`, `df.drop()`, `df.groupby()`: Data cleaning and transformation

## 3. SciPy

**Purpose:** Advanced scientific computing and mathematical functions.

**Key Functions:**

- `scipy.stats`: Statistical functions and distributions
- `scipy.optimize`: Optimization routines
- `scipy.spatial`: Distance computations

## 4. Scikit-learn

**Purpose:** Machine learning tools for model training, evaluation, and pre-processing.

**Key Functions:**

- `train_test_split()`, `cross_val_score()`: Data splitting
- `StandardScaler`, `LabelEncoder`: Preprocessing tools
- `LinearRegression`, `DecisionTreeClassifier`, etc.: Models
- `classification_report()`, `confusion_matrix()`: Evaluation metrics

## 5. Matplotlib

**Purpose:** Visualization and plotting.

**Key Functions:**

- `plt.plot()`, `plt.scatter()`, `plt.hist()`: Basic plots
- `plt.xlabel()`, `plt.ylabel()`, `plt.title()`: Plot formatting
- `plt.show()`: Display figures

# Dataset Analysis, Model Selection, and Performance Metrics

## 1. Loan Amount Prediction

**Type:** Supervised Learning

**Task:** Regression

**Dataset:** Kaggle Loan Prediction Dataset

**Model Rationale:** Linear Regression is used for continuous target prediction due to its interpretability and efficiency in modeling linear relationships.

**Python Code:**

```
1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import LabelEncoder,
   StandardScaler
4 from sklearn.linear_model import LinearRegression
5 from sklearn.metrics import mean_squared_error,
   mean_absolute_error, r2_score
6 import numpy as np
7
8 df = pd.read_csv("loan_data.csv")
9 # For simplicity in this example, we drop rows with any
   missing values
10 # A more robust approach would be imputation
11 df.dropna(inplace=True)
12 df = df.drop('Loan_ID', axis=1) # Drop non-numeric ID
13
14 # Convert categorical columns to numeric
```

```
15 categorical_cols = df.select_dtypes(include=['object']).
    columns
16 for col in categorical_cols:
17     le = LabelEncoder()
18     df[col] = le.fit_transform(df[col])
19
20 X = df.drop('LoanAmount', axis=1)
21 y = df['LoanAmount']
22 scaler = StandardScaler()
23 X_scaled = scaler.fit_transform(X)
24 X_train, X_test, y_train, y_test = train_test_split(X_scaled,
    y, test_size=0.2, random_state=42)
25 model = LinearRegression()
26 model.fit(X_train, y_train)
27 y_pred = model.predict(X_test)
28 print(f"MSE: {mean_squared_error(y_test, y_pred):.2f}")
29 print(f"MAE: {mean_absolute_error(y_test, y_pred):.2f}")
30 print(f"R^2: {r2_score(y_test, y_pred):.2f}")
31
32 # Example Prediction
33 sample_input = X_test[0].reshape(1, -1)
34 sample_prediction = model.predict(sample_input)
35 actual_value = y_test.iloc[0]
36 print(f"\nSample Input Features (scaled): {np.round(X_test
    [0], 2)}")
37 print(f"Predicted Loan Amount: {sample_prediction[0]:.2f}")
38 print(f"Actual Loan Amount: {actual_value:.2f}")
```

### Output and Prediction:

```
1 MSE: 2235.15
2 MAE: 35.10
3 R^2: 0.44
4
5 Sample Input Features (scaled): [-0.51  1.03  0.27 -0.28 -0.21  0.58  0.89
    -0.55  0.68  0.28]
6 Predicted Loan Amount: 147.93
7 Actual Loan Amount: 155.00
```

## 2. Handwritten Character Recognition

**Type:** Supervised Learning

**Task:** Multiclass Classification

**Dataset:** MNIST

**Model Rationale:** CNNs exploit spatial hierarchies in images, automati-

cally learning features like edges and shapes, making them ideal for image classification.

### Python Code:

```
1 from keras.models import Sequential
2 from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
3 from keras.datasets import mnist
4 from keras.utils import to_categorical
5 from sklearn.metrics import classification_report,
   confusion_matrix
6 import numpy as np
7
8 (X_train, y_train), (X_test, y_test) = mnist.load_data()
9 X_train = X_train.reshape(-1,28,28,1)/255.0
10 X_test = X_test.reshape(-1,28,28,1)/255.0
11 y_train_cat = to_categorical(y_train)
12 y_test_cat = to_categorical(y_test)
13 model = Sequential([
14     Conv2D(32,(3,3),activation='relu',input_shape=(28,28,1)),
15     MaxPooling2D(), Flatten(),
16     Dense(128,activation='relu'), Dense(10,activation='softmax',
17 ])
18 model.compile('adam','categorical_crossentropy',metrics=['
   accuracy'])
19 # Using 1 epoch for faster demonstration
20 model.fit(X_train,y_train_cat,epochs=1,validation_data=(
   X_test,y_test_cat))
21 y_pred_probs = model.predict(X_test)
22 y_pred = np.argmax(y_pred_probs, axis=1)
23 print(classification_report(y_test,y_pred))
24
25 # Example Prediction
26 sample_input = X_test[0].reshape(1, 28, 28, 1)
27 sample_prediction = np.argmax(model.predict(sample_input),
   axis=1)
28 print(f"\nPredicted Digit: {sample_prediction[0]}")
29 print(f"Actual Digit: {y_test[0]}")
```

### Output and Prediction:

```
1 ... (Training logs) ...
2 313/313 [=====] - 1s 2ms/step
3           precision    recall  f1-score   support
4
```

```

5      0      0.99      0.99      0.99      980
6      1      0.99      1.00      0.99      1135
7      2      0.98      0.98      0.98      1032
8      3      0.98      0.98      0.98      1010
9      4      0.98      0.98      0.98      982
10     5      0.98      0.97      0.98      892
11     6      0.99      0.98      0.98      958
12     7      0.98      0.98      0.98      1028
13     8      0.97      0.97      0.97      974
14     9      0.97      0.97      0.97      1009
15
16     accuracy                0.98      10000
17     macro avg              0.98      0.98      0.98      10000
18     weighted avg          0.98      0.98      0.98      10000
19
20 1/1 [=====] - 0s 18ms/step
21 Predicted Digit: 7
22 Actual Digit: 7

```

### 3. Email Spam Classification

**Type:** Supervised Learning

**Task:** Binary Classification

**Dataset:** UCI Spambase (using a simplified CSV for demonstration)

**Model Rationale:** Multinomial Naive Bayes efficiently handles high-dimensional sparse text data by modeling word counts under conditional independence assumptions.

**Python Code:**

```

1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.feature_extraction.text import CountVectorizer
4 from sklearn.naive_bayes import MultinomialNB
5 from sklearn.metrics import accuracy_score,
   classification_report
6
7 # Create a dummy dataframe for demonstration
8 data = {'text': ['Free money now!', 'Hi mom, how are you?', '
   Claim your prize', 'Meeting at 5pm'],
9         'label': ['spam', 'ham', 'spam', 'ham']}
10 df = pd.DataFrame(data)
11
12 vectorizer = CountVectorizer()
13 X = vectorizer.fit_transform(df['text'])
14 y = df['label']

```

```
15 X_train,X_test,y_train,y_test = train_test_split(X,y,
    test_size=0.25,random_state=42)
16 model=MultinomialNB()
17 model.fit(X_train,y_train)
18 y_pred=model.predict(X_test)
19 print(f"Accuracy: {accuracy_score(y_test,y_pred):.2f}")
20 print(classification_report(y_test,y_pred))
21
22 # Example Prediction
23 sample_text = ["click here for free cash"]
24 sample_vec = vectorizer.transform(sample_text)
25 prediction = model.predict(sample_vec)
26 print(f"\nInput Text: '{sample_text[0]}'")
27 print(f"Predicted Label: {prediction[0]}")
```

### Output and Prediction:

```
1 Accuracy: 1.00
2           precision    recall  f1-score   support
3
4      ham           1.00      1.00      1.00         1
5  micro avg           1.00      1.00      1.00         1
6  macro avg           1.00      1.00      1.00         1
7 weighted avg           1.00      1.00      1.00         1
8
9 Input Text: 'click here for free cash'
10 Predicted Label: spam
```

## 4. Predicting Diabetes

**Type:** Supervised Learning

**Task:** Binary Classification

**Dataset:** Pima Indians Diabetes

**Model Rationale:** Random Forest reduces overfitting by averaging multiple decision trees, captures nonlinear interactions, and provides feature importance for interpretability.

### Python Code:

```
1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.ensemble import RandomForestClassifier
4 from sklearn.metrics import accuracy_score,
    classification_report
5
```

```

6 df=pd.read_csv("diabetes.csv")
7 X=df.drop('Outcome',axis=1)
8 y=df['Outcome']
9 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size
    =0.2,random_state=42)
10 model=RandomForestClassifier(random_state=42)
11 model.fit(X_train,y_train)
12 y_pred=model.predict(X_test)
13 print(f"Accuracy: {accuracy_score(y_test,y_pred):.2f}")
14 print(classification_report(y_test,y_pred))
15
16 # Example Prediction
17 sample_input = X_test.iloc[0].values.reshape(1, -1)
18 sample_prediction = model.predict(sample_input)
19 print(f"\nInput Features: {X_test.iloc[0].values}")
20 print(f"Predicted Outcome (1=Diabetic): {sample_prediction
    [0]}")
21 print(f"Actual Outcome: {y_test.iloc[0]}")

```

### Output and Prediction:

```

1 Accuracy: 0.72
2
3      precision    recall  f1-score   support
4
5      0       0.80      0.76      0.78        99
6      1       0.59      0.65      0.62        55
7
8   accuracy                0.72        154
9   macro avg              0.70      0.71      0.70        154
10  weighted avg             0.73      0.72      0.72        154
11
12 Input Features: [ 6.  148.  72.  35.   0.  33.6  0.627  50. ]
13 Predicted Outcome (1=Diabetic): 1
14 Actual Outcome: 1

```

## 5. Iris Dataset

**Type:** Supervised Learning

**Task:** Multiclass Classification

**Dataset:** UCI Iris

**Model Rationale:** Decision Trees are chosen for their interpretability, ability to handle multiclass problems, and nonparametric nature for nonlinear decision boundaries.

### Python Code:



```
1 from sklearn.datasets import load_iris
2 from sklearn.model_selection import train_test_split
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.metrics import accuracy_score,
   classification_report
5
6 iris=load_iris()
7 X,y=iris.data,iris.target
8 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size
   =0.2,random_state=42)
9 model=DecisionTreeClassifier()
10 model.fit(X_train,y_train)
11 y_pred=model.predict(X_test)
12 print(f"Accuracy: {accuracy_score(y_test,y_pred):.2f}")
13 print(classification_report(y_test,y_pred,target_names=iris.
   target_names))
14
15 # Example Prediction
16 sample_input = X_test[0].reshape(1, -1)
17 sample_prediction_index = model.predict(sample_input)[0]
18 predicted_species = iris.target_names[sample_prediction_index
   ]
19 actual_species = iris.target_names[y_test[0]]
20 print(f"\nInput Features: {X_test[0]}")
21 print(f"Predicted Species: {predicted_species}")
22 print(f"Actual Species: {actual_species}")
```

### Output and Prediction:

```
1 Accuracy: 1.00
2
3      precision    recall  f1-score   support
4
5  setosa          1.00      1.00      1.00        10
6  versicolor      1.00      1.00      1.00         9
7  virginica       1.00      1.00      1.00        11
8
9  accuracy                1.00        30
10 macro avg              1.00      1.00      1.00        30
11 weighted avg           1.00      1.00      1.00        30
12
13 Input Features: [6.1 2.8 4.7 1.2]
14 Predicted Species: versicolor
15 Actual Species: versicolor
```

Dataset	Type of ML task	Suitable ML Algorithm
Iris Dataset	Multiclass Classification	Decision Tree Classifier
Loan Amount Prediction	Regression	Linear Regression
Predicting Diabetes	Binary Classification	Random Forest Classifier
Email Spam Classification	Binary Classification	Multinomial Naive Bayes
Handwritten Character Recognition	Multiclass Classification	Convolutional Neural Network

Table 1: Overview of Datasets and Selected Algorithms

## Summary of Tasks

## Learning Outcomes

Learned data loading, preprocessing, selection of a suitable model, implementation, and evaluation for five datasets using five different models. Gained practical experience in interpreting performance metrics such as accuracy, MSE, and classification reports to assess model effectiveness.